

AMAN PARGANIHA

+91-6263585818 | aman.parganiha.16@gmail.com | [LinkedIn](#) | [GitHub](#) | [LeetCode](#) | [Portfolio](#)

SUMMARY

AI/ML Engineer (M.Tech, AI/ML) specializing in applied LLM systems, Retrieval-Augmented Generation, and multimodal document AI, with end-to-end MLOps experience from dataset curation and fine-tuning to containerized deployment on AWS. Open-source contributor to Matplotlib and JupyterLab.

EDUCATION

International Institute of Information Technology Naya Raipur India
Master of Technology in Computer Science & Engineering (AI/ML) | CGPA: 8.0/10.0 Aug 2025 – Aug 2027
• Coursework: Advanced Algorithms, Generative AI, Multimodal Systems, Cloud Computing, Distributed Systems, Cybersecurity

Shri Shankaracharya Technical Campus Bhilai India
Bachelor of Technology in Computer Science & Engineering | CGPA: 8.22/10.0 Aug 2019 – Aug 2023
• Coursework: Data Structures & Algorithms, Operating Systems, DBMS, Computer Networks, OOP

TECHNICAL SKILLS

ML & Deep Learning: PyTorch, TensorFlow, Scikit-learn, Transformer architectures, Multimodal models, Fine-tuning, Model Evaluation

Multimodal & Document AI: LlamaIndex, OpenAI API (GPT-4o), RAG pipelines, unstructured data parsing, embedding strategies, OCR concepts

MLOps & Deployment: Docker, AWS, FastAPI, REST APIs, end-to-end ML pipelines, reproducible experimentation, Version Control (Git)

Data Engineering: Pandas, NumPy, ETL pipelines, dataset curation, feature engineering, benchmarking, SQL, NoSQL

Programming: C++ (primary), Python, SQL — clean, modular, production-quality code

Research & Engineering: Paper reading & implementation, NLP, Feature Engineering, Debugging, Unit Testing, Code Reviews

PROJECTS

Event-Driven RAG Agent for Unstructured Document Parsing | [Code](#) | *Python, FastAPI, LlamaIndex, Qdrant, Docker, AWS*
• Built an end-to-end system for parsing and extracting structured information from unstructured PDF documents — directly aligned with Unsiloed AI's core mission of turning complex documents into machine-readable representations.
• Engineered a multimodal retrieval pipeline using LlamaIndex and Qdrant, systematically experimenting with chunk sizes, embedding models, and retrieval strategies — mirrors the benchmarking and optimization work central to this role.
• Integrated GPT-4o with context-aware prompting to generate accurate, grounded answers from indexed technical documents; designed the pipeline to be deterministic and reproducible across document types.
• Containerized the full stack via Docker for end-to-end deployment on AWS; built a production REST API with FastAPI — demonstrates complete MLOps pipeline ownership from model integration to serving.

Multimodal Credit Risk Analysis using SEC XBRL & NLP | [Code](#) | *Python, Pandas, Scikit-learn, NLP, PyTorch, Git*
• Curated a custom multimodal dataset by combining structured SEC XBRL financial filings (2022–2024) with NLP-derived features (sentiment, risk signals, uncertainty metrics) — directly reflects the dataset curation responsibilities of this role.
• Trained, fine-tuned, and rigorously benchmarked binary and multiclass deep learning models, focusing on accuracy, robustness, and efficiency across varied financial document types.
• Implemented modular, reproducible ML pipelines with systematic experimentation tracking — reflecting research-engineer discipline of iterating on models, logging results, and identifying state-of-the-art baselines.

OPEN SOURCE CONTRIBUTIONS

Matplotlib | [PR #31422](#) merged into main (v3.11.0) | Reviewed by core maintainer *timhoffm*
• Improved legend parameter documentation (issue #26620); reviewed and approved by core maintainer timhoffm — demonstrates ability to navigate large open-source codebases and deliver production-quality contributions.

JupyterLab | [PR #18668](#) approved | [PR #18610](#) merged
• Improved user-facing documentation and added Playwright end-to-end test coverage; both PRs approved through multi-round review with core maintainers — reflects research rigor and attention to reproducibility.

CERTIFICATIONS & ACHIEVEMENTS

Google × Kaggle — 5-Day AI Agents Intensive: Hands-on training in agentic LLM workflows, LlamaIndex/LangChain orchestration, and RAG system design.

GATE CSE: Ranked Top 7.8% nationwide (out of ~2.1 lakh candidates) — secured M.Tech (AI/ML) at IIIT Naya Raipur.

Problem Solving: 200+ LeetCode problems solved; consistently top-ranked in programming coursework.